Data Science: Lecture 2

From Concepts to Implementation

Idowu, G.

University of Lagos

What is Data Science?

Definition

"Data Science combines scientific methods, processes, algorithms, and systems to extract knowledge and insights from structured and unstructured data."

Core Components:

- Computer Science
- Statistics
- Domain Knowledge

Key Objectives:

- Descriptive Analytics
- Diagnostic Analytics
- Predictive Analytics
- Prescriptive Analytics

Python Representation

```
#The Data Science Trinity: Computer Science +
    Statistics + Domain Knowledge
```

The Data Science Ecosystem

Computer Science

- Algorithms
- Distributed Computing
- Database Management

Mathematics & Statistics

- Probability Theory
- Linear Algebra
- Statistical Inference

Domain Expertise

- Business Understanding
- Industry Knowledge
- Problem Context

Key Insight

"Without domain knowledge, you're just analyzing numbers without meaning"

CRISP-DM Methodology

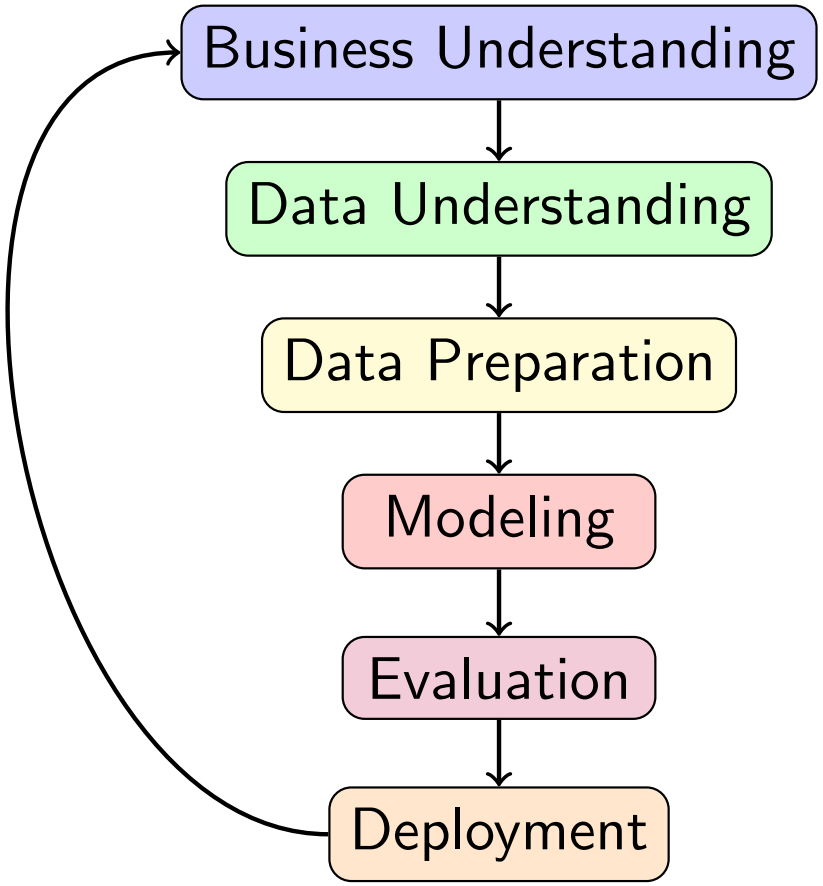

Figure: CRISP-DM Process Flow

CRISP-DM Definition

Cross-Industry Standard Process for Data Mining

Phase 1 - Business Understanding

Key Questions

- What problem are we solving?
- Who are the stakeholders?
- How will we measure success?

Success Criteria Examples

- "Reduce customer churn by 15% in Q3"
- "Achieve 95% accuracy in fraud detection"
- "Decrease inventory costs by 20%"

Common Pitfalls

- Vague problem statements
- Unclear success metrics
- Missing stakeholder alignment

Phase 2 - Data Understanding

Activities:

- Data Collection
- Initial Data Exploration
- Data Quality Assessment
- Pattern Identification

Tools:

- Pandas Profiling
- SweetViz
- Custom EDA scripts

Data Quality Checks

```
# Common data issues
missing_values = data.isnull().sum()
duplicates = data.duplicated().sum()
data_types = data.dtypes
print("Missing values:", missing_values.sum())
print("Duplicate rows:", duplicates)
```

Phase 3 - Data Preparation

Data Cleaning

- Handle missing values
- Remove duplicates
- Correct inconsistencies

Feature Engineering Example

```
# Creating new features
data['age'] = 2024 - data['birth_year']
data['total_spend'] = data['quantity'] * data['price']
data['purchase_month'] = data['purchase_date'].dt.
    month
```

Transformation Techniques

- Normalization (Min-Max Scaling) : Standardization (Z-score)
- One-Hot Encoding : Text Vectorization

Phase 4 - Modeling

Problem Type	Algorithms	Use Cases
Classification	Logistic Regression Random Forest	Spam detection
Regression	Linear Regression Decision Trees	Price prediction
Clustering	K-Means, DBSCAN	Customer segmentation
Dimensionality Reduction	PCA, t-SNE	Data visualization

Table: Algorithm Selection Guide

Model Training Process

1. Train-Test Split (70-30 or 80-20)
2. Model Training on training set
3. Hyperparameter Tuning
4. Cross-validation

Phase 5 - Evaluation

Classification Metrics:

- **Accuracy:** (TP + TN) / Total
- **Precision:** TP / (TP + FP)
- **Recall:** TP / (TP + FN)
- **F1-Score:** Harmonic mean

Regression Metrics:

- MAE
- RMSE
- R^2

Business Validation

- Does the model solve the original problem?
- Is it economically viable?
- Can we explain the results?

Phase 6 - Deployment

Deployment Strategies

- **Batch Processing:** Scheduled predictions
- **Real-time API:** On-demand predictions
- **Embedded Models:** Within applications

Model Serving with Flask

```
from flask import Flask, request, jsonify
import pickle
app = Flask(__name__)
model = pickle.load(open('model.pkl', 'rb'))
@app.route('/predict', methods=['POST'])
def predict():
data = request.get_json()
prediction = model.predict([data['features']])
return jsonify({'prediction': prediction[0]})
```

Machine Learning Types

Supervised Learning

- Uses labeled data
- Predicts known outcomes
- Examples: Classification, Regression

Unsupervised Learning

- Uses unlabeled data
- Discovers hidden patterns
- Examples: Clustering, Association

Reinforcement Learning

- Learns through interaction
- Maximizes cumulative reward
- Examples: Game AI, Robotics

Supervised Learning Deep Dive

Classification Example

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
# Prepare data
X_train, X_test, y_train, y_test = train_test_split(
    features, labels, test_size=0.2)
# Train model
model = RandomForestClassifier(n_estimators=100) model
    .fit(X_train, y_train)
# Predict
predictions = model.predict(X_test)
```

Common Algorithms

Logistic Regression, Decision Trees, Support Vector Machines, Neural Networks

Unsupervised Learning Deep Dive

Clustering with K-Means

```
from sklearn.cluster import KMeans
import matplotlib.pyplot as plt
# Apply K-Means
kmeans = KMeans(n_clusters=3)
clusters = kmeans.fit_predict(customer_data)
# Visualize results
plt.scatter(customer_data['income'], customer_data['
    spending'], c=clusters)
plt.show()
```

Applications

Customer segmentation, Anomaly detection, Market basket analysis, Document clustering

Data Science Toolbox

Programming Languages:

- **Python:** Primary language
- **R:** Statistical computing
- **SQL:** Database querying

Development Environments:

- Jupyter Notebooks
- VS Code
- Google Colab

Python Libraries

```
# Data Manipulation
import pandas as pd
import numpy as np
# Machine Learning
from sklearn import
    ensemble, linear_model
    import tensorflow as tf
# Visualization
import matplotlib.pyplot as
    plt
import seaborn as sns
```

Big Data Technologies

Distributed Computing

- **Apache Spark:** Fast cluster computing
- **Hadoop:** Distributed storage and processing
- **Dask:** Parallel computing in Python

Data Processing Examples

```
# Batch Processing (Spark)
from pyspark.sql import SparkSession
spark = SparkSession.builder.appName(
"BatchProcessing").getOrCreate()
df = spark.read.csv("large_dataset.csv")
# Stream Processing (Kafka + Spark)
stream = spark.readStream.format("kafka").load()
```

Cloud Platforms

AWS SageMaker, Google Cloud AI Platform, Azure Machine Learning

Model Evaluation Techniques

Cross-Validation

```
from sklearn.model_selection import cross_val_score
# 5-fold cross-validation
scores = cross_val_score(model, X, y, cv=5, scoring='
    accuracy')
mean_score = scores.mean()
std_score = scores.std() * 2
print("CV Accuracy:", mean_score, "(+/-", std_score,
    ")")
```

Confusion Matrix

	Positive	Negative
Positive	TP	FN
Negative	FP	TN

ROC Curve

- Trade-off between TP rate and FP rate
- AUC measures overall performance
- Perfect: AUC = 1.0

Feature Engineering Techniques

Numerical Features

- Binning
- Logarithmic transforms
- Polynomial features
- Interaction terms

Categorical Features

- One-Hot Encoding
- Label Encoding
- Target Encoding
- Frequency Encoding

Text Data

- Bag-of-Words
- TF-IDF
- Word Embeddings

Time Series Features

Lag features, Rolling statistics, Seasonal decomposition

Real-World Applications

Healthcare

- Disease prediction
- Medical image analysis
- Drug discovery
- Patient readmission risk

E-commerce

- Recommendation systems
- Price optimization
- Inventory management
- Customer lifetime value

Finance

- Fraud detection
- Algorithmic trading
- Credit scoring
- Customer segmentation

Manufacturing

- Predictive maintenance
- Quality control
- Supply chain optimization
- Energy consumption forecasting

Common Challenges & Solutions

Challenge	Solution
Data Quality Issues	Robust data cleaning pipelines, data validation
Model Performance	Cross-validation, regularization, ensemble methods
Interpretability	SHAP, LIME, feature importance analysis
Deployment	MLOps practices, continuous monitoring
Ethical Concerns	Fairness metrics, differential privacy, ethical guidelines

Table: Common Data Science Challenges and Solutions

The Future of Data Science

Emerging Trends

- **AutoML:** Automated model selection
- **Explainable AI:** Interpretable models
- **MLOps:** DevOps for ML
- **Federated Learning:** Decentralized training
- **AI Ethics:** Responsible AI development

Getting Started in Data Science

Foundation Skills

1. Python programming
2. SQL database querying
3. Statistics and probability
4. Linear algebra

Core Data Science

1. Data manipulation
2. Data visualization
3. Machine Learning
4. Model evaluation

Advanced Topics

1. Deep Learning
2. Big Data technologies
3. Cloud platforms
4. MLOps practices

Practice Platforms

Kaggle competitions, Real-world projects, Open-source contributions

Key Takeaways

1. **Start with the Business Problem**
   - Don't let technology drive the solution
   - Ensure stakeholder alignment
2. **Data Quality is Paramount**
   - Garbage in, garbage out
   - Invest time in data understanding
3. **Iterate and Validate**
   - Use cross-validation
   - Test on unseen data
   - Continuously monitor performance
4. **Think About Deployment Early**
   - Consider production constraints
   - Plan for model maintenance
   - Implement MLOps practices
5. **Communicate Results Effectively**
   - Tell compelling stories with data
   - Visualize insights clearly
   - Focus on business impact

Thank You!

Questions?

"Data science is about extracting knowledge from data, but true value comes from applying that knowledge to drive better decisions."